AR 2EB

ALAMEDA RESEARCH
2: ELECTRIC BOOGALOO
LONG HORIZONS · STRUCTURAL SHIFTS · IMAGINATION

INTERNAL RESEARCH · MEMO · NOT INVESTMENT ADVICE · AI-ASSISTED

Anthropic

PRIVATE : ANTHROPIC · Private Pre-Valuation (exit-scenario)

$61.5B mkt cap · 373M sh · $8B cash, zero debt · Last mark Mar 2025 — Series F; Lead: Lightspeed; Google, Salesforce, Amazon participating

29 May 2026 · the qualitative

$165.00

last round mark · Series F

THE CENTRAL QUESTION

Is the Series F $61.5B mark ($165/share for common) defensible against a probability-weighted distribution of exits over 1-5 years — when secondary buyers are already paying through the last mark in tender offers?

Anthropic prints ~$4B annualized revenue (mid-2026, up from ~$1B end-2024), still growing >100% YoY, with an enterprise mix that increasingly looks like a software platform rather than a research lab. The Series F priced ~15× forward revenue. The dominant variable in valuing the equity today is not an operating DCF (revenue is opaque, inference economics in flux) but the exit distribution: an IPO between 2027 and 2029 at $250B-$1.5T is the modal good outcome. The bear case is not slow erosion but a capability eclipse — OpenAI ships a step-function, AWS cools on Bedrock, talent exits to Google DeepMind. Probability-weighted present value lands near 2× the last mark, justifying secondary buyers paying premium-to-mark. The single most consequential editorial input is the discount rate, not the exit valuations. Five exit scenarios, weighted; show your work.

PROBABILITY-WEIGHTED EXPECTED VALUE

$346.69 expected fair value today

+110.1% vs last mark $165.00

WEIGHTING RATIONALE

Wipeout 10% Capability eclipse + AWS withdrawal + talent flight — compound 3-vector tail.
Slow Private 20% Stripe analog: indefinitely private, LTBT-blessed, holders illiquid.
Base IPO 42% Modal: 2027 IPO at ~$250B, ~4× last-round; routine cohort dynamics.
Bull IPO 20% Capability lead holds 24+ months, enterprise NDR sustains, $700B IPO 2028.
Ultra Bull 8% AGI-tier moat + regulatory positioning; $1.5T tier-1 mega-cap club.

Bull IPO (20%) + Ultra Bull (8%) contribute $225 — 65% of the $347 weighted PV despite 28% combined probability. The last round mark ($165) sits 52% below the weighted PV, so secondary buyers paying through the mark are roughly justified. The venture discount rate is the load-bearing input — see Page 4 sensitivity.

FUNDING ROUNDS · LAST MARK · EXIT-SCENARIO FAN (PV at exit, $/sh)

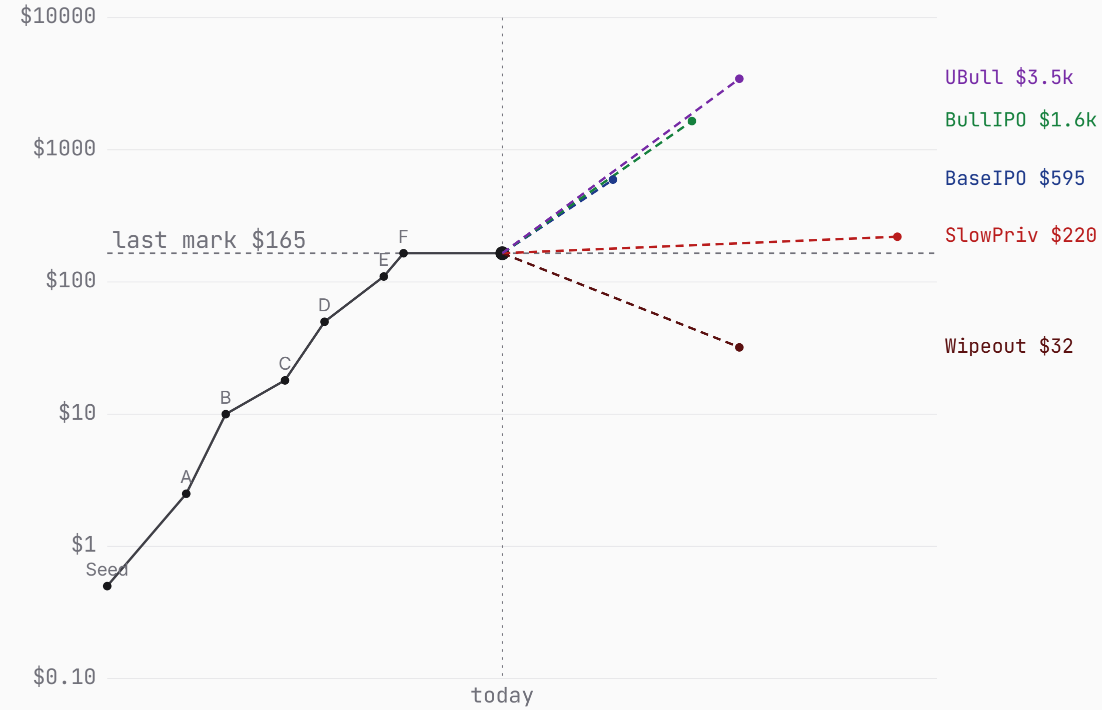

WIPEOUT	SLOW PRIVATE	BASE IPO	BULL IPO	ULTRA BULL
$14.10	$78.14	$250.13	$543.45	$1448.85
-91.5% vs spot	-52.6% vs spot	+51.6% vs spot	+229.4% vs spot	+778.1% vs spot
Exit $12B Yrs 3.0 WACC 22% Prob 10%	Exit $95B Yrs 5.0 WACC 23% Prob 20%	Exit $250B Yrs 1.4 WACC 18% Prob 42%	Exit $700B Yrs 2.4 WACC 18% Prob 20%	Exit $1.5T Yrs 3.0 WACC 18% Prob 8%
Capability eclipsed; down-round.	Stays private 5y+; no liquidity.	2027 IPO at $250B; 1.5× the mark.	2028 IPO at $700B; lead holds.	AGI-tier moat; $1.5T IPO.

AR 2EB | ALAMEDA RESEARCH 2: ELECTRIC BOOGALOO · LONG HORIZONS · STRUCTURAL SHIFTS · IMAGINATION

Anthropic scenario narratives

PROBABILITY WEIGHTS

Ultra Bear 10% Bear 20% Base 42% Bull 20% Ultra Bull 8% · Weighted expected $346.69 (+110.1% vs spot)

WIPEOUT $14.10

-91.5% vs spot

Probability 10%

Capability eclipsed; down-round.

WHY 10%

Capability eclipse + safety-talent attrition + AWS withdrawal is a compound 3-vector event. Each individually 25-40% plausible across a 3-year horizon; joint at ~10%. The dominant risk for any frontier-AI lab is being lapped on capability — the downside is one binary, not slow erosion. Bayesian update from 2026: Claude 4.5 reception was strong, enterprise revenue held, the RSP framework didn't constrain commercial deployment as much as bears feared — so a-priori 15% nudges down to 10%.

WHAT HAPPENS

OpenAI ships GPT-6 with a step-function capability gain. Google folds Gemini-Ultra into Workspace at near-zero marginal price. Anthropic's growth decelerates; AWS cools on Bedrock as Amazon's own Nova line gains traction. Top safety researchers exit for Google DeepMind and a stealth lab. A planned 2027 IPO gets pulled; a 2028 bridge round prices below the Series F mark, triggering dilution and a quiet acqui-hire by Amazon or Google.

At the resulting $12B post-money, common-equity per share falls to ~$32 nominal — discounted three years at a 22% venture-WACC (illiquidity + execution + dilution overhang), and after a 20% within-scenario write-to-zero probability, expected PV is ~$14/share, roughly a tenth of the last mark.

SLOW PRIVATE $78.14

-52.6% vs spot

Probability 20%

Stays private 5y+; no liquidity.

WHY 20%

Stripe is the canonical analog: $95B mark, no IPO seven years later. Anthropic's PBC + Long-Term Benefit Trust structure makes the "we don't have to go public" stance more credible than for a vanilla C-corp. Modal over 3 years but not over 5+; at some point holders demand liquidity. 20% reflects the real possibility the LTBT prefers indefinite private status.

WHAT HAPPENS

The "stays private indefinitely" path. Anthropic raises 2-3 more growth rounds at marks that grow with the AI sector — Series G ~2027 around $95B, Series H ~2029 around $130B — but no IPO appears credible (regulatory uncertainty, a depressed IPO window, or a Long-Term Benefit Trust decision that public-company quarterly pressure conflicts with the safety mission). Secondaries create episodic liquidity at steep discounts to the mark.

The economics are real — ARR grows from $4B to ~$25B by 2031 — but without an IPO catalyst, the holder faces ~5 years of illiquidity and substantial dilution risk. PV at a 23% venture discount: ~$78, meaningfully below the $165 last mark even as the nominal round price grows. This is the underappreciated outcome.

BASE IPO $250.13

+51.6% vs spot

Probability 42%

2027 IPO at $250B; 1.5× the mark.

WHY 42%

Modal outcome. Conditions: enterprise revenue compounds >100% YoY through end-2026; no catastrophic safety incident delays timing; the AI IPO window opens on the back of OpenAI's secondary or its own IPO. Each individually ~70-80%; joint ~40-45%. 42% reflects that this IS the base case but isn't overwhelming — Anthropic could slide into slow_private or bull_ipo depending on conditions.

WHAT HAPPENS

Anthropic IPOs mid-2027 on the heels of an enterprise ARR inflection ($8B+ run-rate by Q2 2027), riding a re-opened AI IPO window that OpenAI's late-2026 secondary tender priced. Lead underwriters anchor a $250B post-money — roughly 4× the Series F valuation — reflecting both fundamental progress (Claude family revenue, Bedrock and Vertex pass-through) and AI-sector multiple expansion. Constitutional AI / RSP positioning becomes a marketed governance differentiator vs OpenAI.

Per-share at IPO ≈ $595. PV back to today at an 18% venture-WACC over ~1.4 years = ~$472 — but with a within-scenario haircut for IPO underperformance (47%, reflecting lockup decay + typical post-lockup behavior across the 2027 AI cohort), expected PV/share is ~$250. That is 1.5× the $165 last mark — explaining why secondary buyers pay above-mark in tender offers running into the IPO window.

BULL IPO $543.45

+229.4% vs spot

Probability 20%

2028 IPO at $700B; lead holds.

WHY 20%

Requires conjunction: capability lead persists 24+ months; enterprise NDR holds >130%; gross margin expands as inference compute commoditizes downward; the 2028 AI IPO window receives the cohort favorably. Each 40-60% individually; joint ~20%. Higher floor than a typical young-company bull because Anthropic already has revenue traction.

WHAT HAPPENS

Anthropic's capability lead consolidates. Claude 5 / 6 deliver durable enterprise stickiness — net retention >150%, gross-margin expansion through caching + lower inference unit costs. The AWS Bedrock relationship deepens; Vertex pass-through grows. Enterprise share of revenue rises above 80%, making the business look more like Snowflake or Databricks than consumer-AI. Late-2028 IPO at $700B reflects ARR ($18-22B run-rate) and multiple expansion to 35-40× ARR.

Per-share at IPO ≈ $1,650 nominal; PV at 18% over 2.4 years = ~$1,109; the IPO-underperformance haircut brings expected PV/share to ~$543. A 3.3× outcome on the last mark — material upside that justifies the secondary buyer paying through the mark today.

ULTRA BULL $1448.85

+778.1% vs spot

Probability 8%

AGI-tier moat; $1.5T IPO.

WHY 8%

Tail of tails. Each Bull condition AND a capability gap that justifies mega-cap multiples AND regulatory positioning as the "safe lab" AND a durable lead over OpenAI / Google / xAI for 3+ years. 5-10% feels right; 8% acknowledges Anthropic's safety-research depth without over-weighting the tail.

WHAT HAPPENS

The genuine tail. Anthropic produces an AGI-tier capability gap that monetizes as a software platform — autonomous-agent revenue compounds, Claude becomes the default enterprise AI substrate, and revenue inflects from $20B (FY28) to $50B+. Constitutional AI / RSP becomes a regulatory moat (Anthropic is "the safe lab" in the EU AI Act and US AI policy regime). IPO 2029 at $1.5T — joining Microsoft / Google / Meta in the tier-1 mega-cap club.

Per-share at IPO ≈ $3,450 nominal; PV at 18% over 3 years = ~$2,100; the haircut brings expected PV/share to ~$1,449 — an 8.8× outcome on the last mark. Contributes the largest single slice of weighted value despite only 8% probability — compound-conditional-tail behavior.

INTERNAL RESEARCH · MEMO · NOT INVESTMENT ADVICE · AI-ASSISTED

Anthropic business snapshot

29 May 2026 · the snapshot

Bear $78.14 Base $250.13 Bull $543.45

Funding-round history + exit-scenario distribution · no operating-business DCF (revenue opaque) · discount rate is the load-bearing input

ARR build to scenario exit (run-rate $B)

Exit valuation distribution (post-money $B)

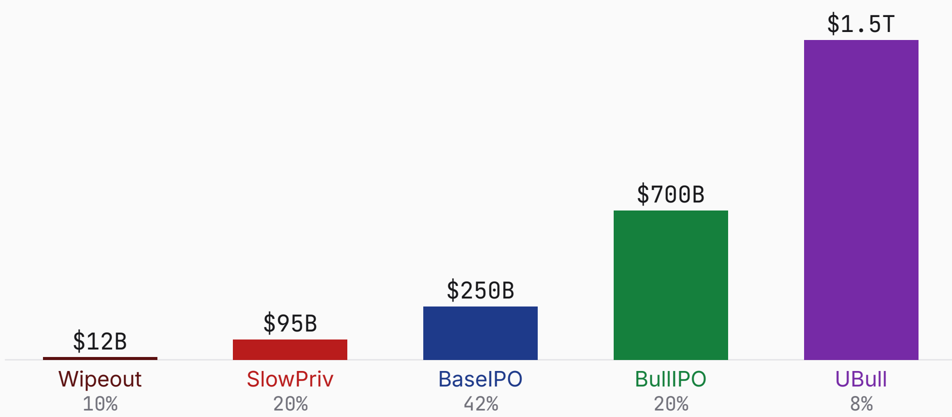

Sources: TechCrunch, Bloomberg, The Information, Anthropic blog, AWS announcements, SEC filings of Amazon/Google equity disclosures.

Anthropic 7 Powers · the moat, scored

Bear $78.14 Base $250.13 Bull $543.45

Arena. Frontier AI labs — Anthropic vs OpenAI, Google DeepMind, xAI, Meta AI; the race for frontier capability plus the compute, capital, and talent to sustain it.

POWER ORIGINATION
origination window: open

POWER ORIGINATION · 7 POWERS

Scale Economies	Training-cost amortization needs scale; well-capitalized (Amazon/Google) but smaller than OpenAI/Google.
Network Economies	Claude developer ecosystem + MCP standard emerging; far less consumer network than ChatGPT.
Counter-Positioning	Safety-first, enterprise-trust positioning vs consumer-first incumbents (Constitutional AI).
Switching Costs	Enterprise API integration; Claude embedded in coding workflows; MCP.
Branding	Strong in enterprise + coding + safety; less consumer mindshare than ChatGPT/Gemini.
Cornered Resource · thesis leans here	Elite research talent; Amazon (~$8B) + Google compute/Trainium access.
Process Power	Research velocity; RL + interpretability/safety methods.

THE RACE · NAMED RIVALS

OpenAI — $ tens of B + Microsoft
~32% terminal share · ChatGPT consumer scale + Microsoft; largest war chest.

Google DeepMind — Alphabet balance sheet
~24% terminal share · Gemini + TPUs + distribution + balance sheet; the scale threat.

xAI — $ billions raised
~12% terminal share · Grok + X data + Musk capital + large GPU cluster.

Meta AI — Meta balance sheet
~12% terminal share · Open-weight Llama; free distribution at platform scale.

LEAD / LAG · FALSIFIERS

Frontier capability (coding/agents)	leading vs parity	LEADING
Compute / capital scale	~$8-13B vs far larger	LAGGING
Consumer distribution	limited vs massive	LAGGING

> **COMPETITIVE READ**
> Anthropic is originating a real enterprise + coding + safety position with elite talent and major-cloud compute, but it lags OpenAI and Google on capital scale and consumer reach; the window is open while the frontier moves — the compute/capital gap is the falsifier for whether it stays at the frontier.

Anthropic show your work

Bear $78.14 Base $250.13 Bull $543.45 · Weighted $346.69 (+110.1%)

WIPEOUT $14.10

ASSUMPTIONS

Exit kind	down round
Exit year	2029
Exit post-money	$12B
Per-share at exit ($)	32
Years to exit	3.0
Venture WACC (%)	22
Dilution to exit (%)	+14
P(total loss) (%)	20
Scenario probability (%)	10

EXIT → PV PER SHARE

Per-share at exit (nominal)	$32
÷ (1+22%)^3.0	$17.62
× (1−P_loss 20%)	
= Expected PV per share	**$14.10**

PV SENSITIVITY TO DISCOUNT RATE

15%	18%	22%	25%
$17	$16	$14	$13

SLOW PRIVATE $78.14

ASSUMPTIONS

Exit kind	stay private
Exit year	2031
Exit post-money	$95B
Per-share at exit ($)	220
Years to exit	5.0
Venture WACC (%)	23
Dilution to exit (%)	+16
P(total loss) (%)	0
Scenario probability (%)	20

EXIT → PV PER SHARE

Per-share at exit (nominal)	$220
÷ (1+23%)^5.0	$78.14
= Expected PV per share	**$78.14**

PV SENSITIVITY TO DISCOUNT RATE

15%	18%	22%	25%
$109	$96	$81	$72

BASE IPO $250.13

ASSUMPTIONS

Exit kind	ipo
Exit year	2027
Exit post-money	$250B
Per-share at exit ($)	595
Years to exit	1.4
Venture WACC (%)	18
Dilution to exit (%)	+12
P(total loss) (%)	0
Scenario probability (%)	42

EXIT → PV PER SHARE

Per-share at exit (nominal)	$595
÷ (1+18%)^1.4	$471.93
× (1−IPO haircut 47%)	$250.12
= Expected PV per share	**$250.13**

PV SENSITIVITY TO DISCOUNT RATE

15%	18%	22%	25%
$259	$250	$239	$231

BULL IPO $543.45

ASSUMPTIONS

Exit kind	ipo
Exit year	2028
Exit post-money	$700B
Per-share at exit ($)	1,650
Years to exit	2.4
Venture WACC (%)	18
Dilution to exit (%)	+14
P(total loss) (%)	0
Scenario probability (%)	20

EXIT → PV PER SHARE

Per-share at exit (nominal)	$1,650
÷ (1+18%)^2.4	$1109.09
× (1−IPO haircut 51%)	$543.45
= Expected PV per share	**$543.45**

PV SENSITIVITY TO DISCOUNT RATE

15%	18%	22%	25%
$578	$543	$502	$473

ULTRA BULL $1448.85

ASSUMPTIONS

Exit kind	ipo
Exit year	2029
Exit post-money	$1.50T
Per-share at exit ($)	3,450
Years to exit	3.0
Venture WACC (%)	18
Dilution to exit (%)	+17
P(total loss) (%)	0
Scenario probability (%)	8

EXIT → PV PER SHARE

Per-share at exit (nominal)	$3,450
÷ (1+18%)^3.0	$2099.78
× (1−IPO haircut 31%)	$1448.85
= Expected PV per share	**$1448.85**

PV SENSITIVITY TO DISCOUNT RATE

15%	18%	22%	25%
$1565	$1449	$1311	$1219

Anthropic supporting analysis · disclaimers · glossary

Bear $78.14 Base $250.13 Bull $543.45

PUSHBACK · WHY THE BASE CASE IS TOO HARSH

1 **Enterprise revenue inflection is real.**
$1B → $4B in 18 months beats Snowflake's comparable phase. NDR >150% implies land-and-expand, not pilots.

2 **AWS Project Rainier is a strategic moat.**
A 1M+ Trainium2 cluster reserved for Anthropic is durable compute alignment rivals can't replicate cheaply.

3 **Safety positioning may monetize in regulation.**
EU AI Act enforcement begins 2026-2027. RSP / Constitutional AI framing makes Anthropic the easiest model to deploy in regulated industries.

4 **Secondaries pay through the mark.**
2026 tender activity cleared ~$200/share-equiv (~$74B implied) for common — the equity already trades above the $165 mark.

FALSIFICATION TRIGGERS

Bear validation

OpenAI ships GPT-6 with measurable step-function gains · 200+ research-head exit in any 6-month window · AWS cools publicly on Project Rainier expansion

Bull validation

ARR crosses $8B run-rate by Q2 2027 · a confidential S-1 leaks · enterprise NDR sustains >130% in a published reference disclosure

Reframe needed

An LTBT action blocking/delaying IPO — the 'stay private' path becomes modal, inverting the discount-rate logic and depressing PV ~40%.

DISCLAIMERS · PLEASE READ BEFORE USING THIS DOCUMENT

Not investment advice.

This document is produced for entertainment and educational purposes only by an individual amateur investor. "Alameda Research 2: Electric Boogaloo" is the unincorporated personal concept of one person and is not a registered investment management entity. Nothing herein constitutes investment advice, a recommendation, or a solicitation to buy or sell any security.

Not a professional.

The author is not a registered investment advisor, broker-dealer, CFA charterholder, certified financial planner, or licensed financial professional of any kind. The author has no formal training in equity research or capital markets. No fiduciary relationship is created by reading this document.

AI-assisted analysis.

Substantial portions of this analysis were produced with the assistance of large language models. LLMs are known to fabricate facts, miscalculate numbers, hallucinate sources, and present incorrect information with high apparent confidence. Every figure, claim, and projection in this document should be independently verified before acting on it.

Forward-looking statements.

Scenarios, valuations, projections, and any forward-looking statements involve substantial assumptions and uncertainty. Actual results may differ materially from those projected. Past performance is not indicative of future results. The model is a model; reality is not.

Author may hold positions.

The author may own, intend to acquire, or hold short exposure to $ANTHROPIC or related securities at any time. Positions may change without notice and without an update to this document. Do not assume the author's positions match the tone of the analysis.

Do your own research.

Do not make investment decisions on the basis of this document. Consult a licensed financial professional and read the company's official SEC filings (10-K, 10-Q, 8-K, proxy) before forming any investment view. If you wouldn't trust your retirement to a chatbot, don't trust this either.

GLOSSARY · CONCEPTS REFERENCED IN THE NARRATIVE

Constitutional AI (CAI) — Anthropic's training method — models trained against a written constitution using AI feedback rather than dense human ratings.

Responsible Scaling Policy (RSP) — Published commitments to halt/slow deployments when dangerous-capability evals trip. Both a safety mechanism and a marketed differentiator.

Claude family (Opus / Sonnet / Haiku) — Tiered model family — Opus highest capability, Sonnet the workhorse, Haiku fast/cheap. Generations refresh ~twice a year.

Long-Term Benefit Trust (LTBT) — Independent trust holding a stock class that elects a board majority over time. Non-investor members aligned with the safety mission.

Project Rainier — AWS's announced 1M+ Trainium2 cluster reserved for Anthropic — a supply commitment and a marker of AWS's strategic stake.

Last-round mark — Per-share value implied by the latest round (post-money ÷ fully-diluted). A preferred, illiquid, single-counterparty price — not a market quote.